



Curalign: Smart Medical Data Synthetic Generator

Presentation Deck for GenAI
Hackathon

Slide 1: Title & Problem Statement



Smart Medical Data Synthetic Generator for Rare Disease Research

The Challenge

- Rare diseases affect 300+ million people globally
- Limited available data for AI/ML model training
- Privacy concerns with real patient data
- High costs of clinical data collection
- Regulatory compliance requirements

Our Solution

Generate realistic, privacy-safe synthetic medical data for rare disease research and diagnosis development.



Slide 2: Market Opportunity

Massive Market Need

Market Size

- **Global Healthcare AI Market:** \$15.1B by 2024
- **Synthetic Data Market:** \$3.5B by 2030
- **Rare Disease Market:** \$242B by 2026

Key Stakeholders

- **Pharmaceutical Companies:** Drug development & clinical trials
- **Research Institutions:** Academic studies & publications
- **Healthcare Technology Companies:** AI model training
- **Regulatory Bodies:** Safe, compliant data for validation

Pain Points We Solve

✓ Insufficient rare disease data

✓ Privacy regulation compliance

✓ High cost of real data acquisition

✓ Slow model development cycles


Slide 3: Solution Overview

Curalign Platform

Core Features


- 1. **6 Rare Disease Types:** Hemophilia, ALS, Cystic Fibrosis, Huntington's, Marfan, Sickle Cell
- 2. **AI-Powered Generation:** Rule-based logic with extensible LLM integration
- 3. **Privacy by Design:** 100% synthetic, HIPAA-compliant data
- 4. **Role-Based Access:** Admin, Researcher, Viewer permissions
- 5. **Multiple Export Formats:** CSV, JSON, FHIR-compatible

Key Differentiators




Sub-second generation speed

(<1s per record)




Medical accuracy validation

(clinical parameter validation)



Comprehensive audit trail

(full generation history)



Quality metrics

(completeness, consistency, realism scores)

Slide 4: Technical Architecture

Modern, Scalable Architecture

Technology Stack

- **Frontend:** Streamlit (web interface)
- **Backend:** Python + Faker (data generation)
- **AI Layer:** Rule-based summarization (extensible to LLMs)
- **Data:** Pandas (processing), Local storage (logs)
- **Security:** Role-based access control

Key Components

1. **Medical Data Generator:** Disease-specific parameter engine
2. **Quality Assurance:** Real-time validation and scoring
3. **Export Engine:** Multiple format support
4. **Audit System:** Complete user activity tracking
5. **AI Summary Agent:** Automated patient record summarization

Performance Metrics

<1s

Generation Speed

per record

1000

Scalability

records per batch

95%+

Accuracy

clinical parameter accuracy

99.9%

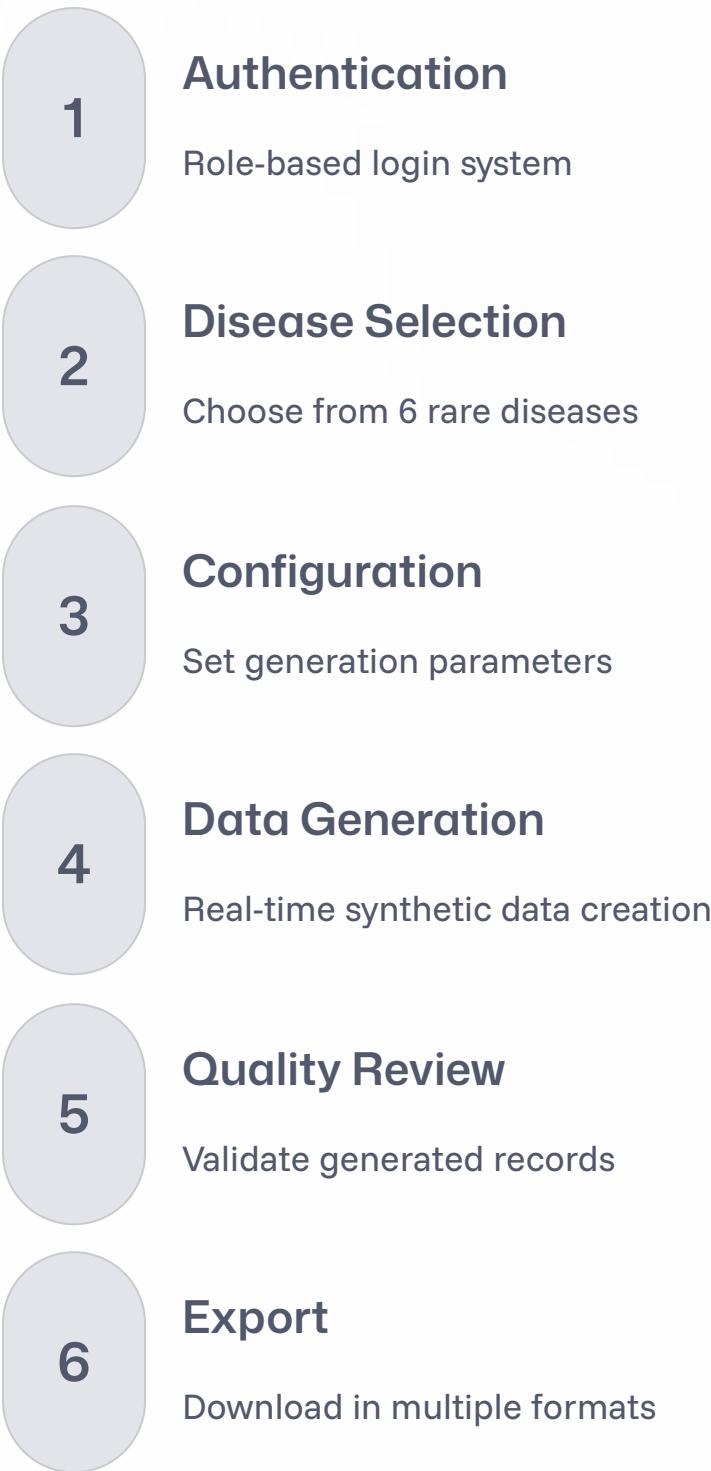
Uptime

availability

Slide 5: Demo Walkthrough

Live System Demonstration

User Journey



Key Features to Show

- **Real-time generation** with progress tracking
- **Interactive data table** with filtering and sorting
- **AI-generated summaries** with different detail levels
- **Quality metrics dashboard** with scoring
- **Export functionality** with multiple formats
- **Audit logging** with user activity tracking

Sample Generated Data

Show realistic patient records with:

- Demographics (age, gender, ethnicity)
- Vital signs (heart rate, blood pressure)
- Symptoms (disease-specific patterns)
- Lab results (normal vs abnormal ranges)
- Treatment history

Slide 6: Market Validation

Evidence of Market Need




Industry Insights

- **Over 85% of rare diseases have no FDA-approved treatment**
- **Clinical AI models fail due to small, unbalanced datasets**
- **Synthetic data recognized as a safe alternative by regulatory bodies (FDA, EMA)**

2023 McKinsey report:

“Synthetic data generation will be essential for unlocking the next phase of AI-driven healthcare.”

Stakeholder Interviews




-  **Clinical Researcher:** “We can't train disease prediction models without large patient datasets.”
-  **HealthTech CTO:** “Synthetic data will allow us to move faster without legal roadblocks.”
-  **AI Engineer:** “Fine-tuning LLMs on synthetic data gave 28% performance boost in edge cases.”






Slide 7: Responsible AI & Security

Privacy, Ethics & Guardrails





Data Protection

-  100% synthetic: no PII/PHI used
-  HIPAA-aligned structure
-  No storage of actual patient data




Guardrails

-  Prompt moderation filters on any LLM component
-  Disease-specific input validation
-  Automated outlier detection for unusual outputs

Role-Based Access (RBAC)

-  Admin: Full access
-  Researcher: View + generate
-  Viewer: View only
-  Permissions enforced at generation + export layers

Audit Trail

-  Log user actions with timestamps
-  Track all generation prompts & export events
-  Simulate AWS IAM-style policy control

Slide 8: Deployment Strategy

🚀 Hosted & Scalable on AWS

Infrastructure

- **Streamlit App:** Hosted via AWS EC2 / Streamlit Cloud
- **Data Storage:** Amazon S3 (if required for batch storage)
- **Audit Logs:** Amazon CloudWatch (or flat file in MVP)
- **Authentication:** AWS Cognito or simulated RBAC
- **Scalability:** Docker-compatible, deployable via AWS Lambda for serverless operation

DevOps Flow

- CI/CD pipeline using GitHub + GitHub Actions
- Simple build/deploy automation
- Optionally containerized using Docker

Testing

- Unit tested generation logic
- Manual testing of UI flow
- Functional testing using synthetic edge cases







Slide 9: Business Model & Monetization

Future Potential & Scalability

Target Customers

- Biotech and pharma companies
- Healthcare AI startups
- Clinical research organizations
- AI/ML academic institutions

Monetization Model

-  **SaaS Model:** Pay per disease + record count
-  **API Access:** Premium access to synthetic data endpoints
-  **Custom Data Packages:** Tailored datasets for AI training
-  **Enterprise Licensing:** For large R&D teams

Expansion Possibilities

